

Agent-based modeling as generative explanation

Growing macro-patterns from micro-level cognitive rules

In the previous lecture, the core problem was workflow control:

In the previous lecture, the core problem was workflow control:

- define the goal

In the previous lecture, the core problem was workflow control:

- define the goal
- constrain the implementation

In the previous lecture, the core problem was workflow control:

- define the goal
- constrain the implementation
- verify the result

In the previous lecture, the core problem was workflow control:

- define the goal
- constrain the implementation
- verify the result

In the previous lecture, the core problem was workflow control:

- define the goal
- constrain the implementation
- verify the result

ABM gives us a modeling problem where that discipline matters immediately.

Cognitive models describe how one person:

Cognitive models describe how one person:

- forms beliefs

Cognitive models describe how one person:

- forms beliefs
- updates on evidence

Cognitive models describe how one person:

- forms beliefs
- updates on evidence
- makes decisions under uncertainty

Cognitive models describe how one person:

- forms beliefs
- updates on evidence
- makes decisions under uncertainty

Cognitive models describe how one person:

- forms beliefs
- updates on evidence
- makes decisions under uncertainty

This is one standard level of analysis in cognitive science.

Many phenomena of interest are not properties of a single mind.

Many phenomena of interest are not properties of a single mind.

- Why do populations polarize even when individuals reason well?

Many phenomena of interest are not properties of a single mind.

- Why do populations polarize even when individuals reason well?
- Why do norms emerge and persist without central enforcement?

Many phenomena of interest are not properties of a single mind.

- Why do populations polarize even when individuals reason well?
- Why do norms emerge and persist without central enforcement?
- Why do crowds sometimes find good solutions and sometimes fail?

Many phenomena of interest are not properties of a single mind.

- Why do populations polarize even when individuals reason well?
- Why do norms emerge and persist without central enforcement?
- Why do crowds sometimes find good solutions and sometimes fail?

Many phenomena of interest are not properties of a single mind.

- Why do populations polarize even when individuals reason well?
- Why do norms emerge and persist without central enforcement?
- Why do crowds sometimes find good solutions and sometimes fail?

A single-agent model can specify the cognitive mechanism, but it cannot by itself explain the population-level pattern.

How can population-level polarization emerge even when individuals only make small, locally reasonable belief updates?

An illustrative question

How can population-level polarization emerge even when individuals only make small, locally reasonable belief updates?

We will treat this as a generative question, not a correlational one.

- A standard cognitive model asks how one person updates.

- A standard cognitive model asks how one person updates.
- An ABM asks what many interacting updaters produce together.

- A standard cognitive model asks how one person updates.
- An ABM asks what many interacting updaters produce together.
- ABM embeds cognitive mechanisms inside an interaction structure.

What is an agent?

An agent is a simplified decision-maker with internal state, local inputs, and update rules.

What is an agent?

An agent is a simplified decision-maker with internal state, local inputs, and update rules.

- State changes over time.

What is an agent?

An agent is a simplified decision-maker with internal state, local inputs, and update rules.

- State changes over time.
- Rules map local inputs to local actions.

- belief

Agent state examples

- belief
- memory

Agent state examples

- belief
- memory
- attention

Agent state examples

- belief
- memory
- attention
- confidence

Agent state examples

- belief
- memory
- attention
- confidence
- strategy

Agent state examples

- belief
- memory
- attention
- confidence
- strategy
- preference

Agent state examples

- belief
- memory
- attention
- confidence
- strategy
- preference
- location

Agent state examples

- belief
- memory
- attention
- confidence
- strategy
- preference
- location
- social ties

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

- agents and their states

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

- agents and their states
- update rules

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

- agents and their states
- update rules
- environment and interaction structure

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

- agents and their states
- update rules
- environment and interaction structure
- update schedule

What is an agent-based model?

An agent-based model is a computational model in which possibly heterogeneous agents update their states through local rules while interacting with each other or with an environment.

Components:

- agents and their states
- update rules
- environment and interaction structure
- update schedule
- macro-level outcomes

- Schelling (1971): residential segregation from mild local preference

- Schelling (1971): residential segregation from mild local preference
- Axelrod (1984): cooperation from iterated local interaction

- Schelling (1971): residential segregation from mild local preference
- Axelrod (1984): cooperation from iterated local interaction
- Epstein and Axtell (1996): **Growing Artificial Societies**

- Schelling (1971): residential segregation from mild local preference
- Axelrod (1984): cooperation from iterated local interaction
- Epstein and Axtell (1996): **Growing Artificial Societies**
- Epstein (1999): “If you didn’t grow it, you didn’t explain it”

- Schelling (1971): residential segregation from mild local preference
- Axelrod (1984): cooperation from iterated local interaction
- Epstein and Axtell (1996): **Growing Artificial Societies**
- Epstein (1999): “If you didn’t grow it, you didn’t explain it”

- Schelling (1971): residential segregation from mild local preference
- Axelrod (1984): cooperation from iterated local interaction
- Epstein and Axtell (1996): **Growing Artificial Societies**
- Epstein (1999): “If you didn’t grow it, you didn’t explain it”

In cognitive science, ABM connects naturally to social learning, cultural evolution, and language emergence.

To explain a macro-pattern, grow it from micro-level rules.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.
3. Run the system.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.
3. Run the system.
4. Check whether the target pattern appears.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.
3. Run the system.
4. Check whether the target pattern appears.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.
3. Run the system.
4. Check whether the target pattern appears.

The explanatory claim is mechanism-level sufficiency.

To explain a macro-pattern, grow it from micro-level rules.

1. Specify agents and their states.
2. Specify interaction and update rules.
3. Run the system.
4. Check whether the target pattern appears.

The explanatory claim is mechanism-level sufficiency.

But ABM does not license simple inverse inference: the same macro-pattern can arise from different micro-mechanisms.

A pattern is **emergent** when it appears at the population level even though no individual agent represents, intends, or directly controls that pattern.

A pattern is **emergent** when it appears at the population level even though no individual agent represents, intends, or directly controls that pattern.

- Each agent: a simple belief-update rule, nothing more.

A pattern is **emergent** when it appears at the population level even though no individual agent represents, intends, or directly controls that pattern.

- Each agent: a simple belief-update rule, nothing more.
- Population: a bimodal distribution of beliefs develops.

A pattern is **emergent** when it appears at the population level even though no individual agent represents, intends, or directly controls that pattern.

- Each agent: a simple belief-update rule, nothing more.
- Population: a bimodal distribution of beliefs develops.

A pattern is **emergent** when it appears at the population level even though no individual agent represents, intends, or directly controls that pattern.

- Each agent: a simple belief-update rule, nothing more.
- Population: a bimodal distribution of beliefs develops.

Emergence is the central explanandum of ABM.

- Statistical model: which variables predict observed outcomes?

- Statistical model: which variables predict observed outcomes?
- Cognitive model: what mechanism describes one agent?

ABM and statistical models answer different questions

- Statistical model: which variables predict observed outcomes?
- Cognitive model: what mechanism describes one agent?
- ABM: what happens when many such mechanisms interact?

ABM and statistical models answer different questions

- Statistical model: which variables predict observed outcomes?
- Cognitive model: what mechanism describes one agent?
- ABM: what happens when many such mechanisms interact?

ABM and statistical models answer different questions

- Statistical model: which variables predict observed outcomes?
- Cognitive model: what mechanism describes one agent?
- ABM: what happens when many such mechanisms interact?

Good practice triangulates across all three.

- cultural transmission

Why cognitive scientists should care

- cultural transmission
- social learning

Why cognitive scientists should care

- cultural transmission
- social learning
- collective intelligence

Why cognitive scientists should care

- cultural transmission
- social learning
- collective intelligence
- polarization

Why cognitive scientists should care

- cultural transmission
- social learning
- collective intelligence
- polarization
- language emergence

Why cognitive scientists should care

- cultural transmission
- social learning
- collective intelligence
- polarization
- language emergence
- developmental systems

Why cognitive scientists should care

- cultural transmission
- social learning
- collective intelligence
- polarization
- language emergence
- developmental systems
- cognitive ecology

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning
- categorization

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning
- categorization
- decision thresholds

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning
- categorization
- decision thresholds
- confidence-weighted learning

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning
- categorization
- decision thresholds
- confidence-weighted learning

What makes an ABM cognitive?

An ABM becomes a cognitive-science model when the agent rule represents a theory of cognition.

- belief updating
- memory retrieval
- attention allocation
- reinforcement learning
- categorization
- decision thresholds
- confidence-weighted learning

The population model is only as cognitive as the agent model.

Toy model or empirical model?

Toy ABM:

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

- constrains rules from data

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

- constrains rules from data
- compares simulated and observed patterns

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

- constrains rules from data
- compares simulated and observed patterns
- treats validation as part of the model

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

- constrains rules from data
- compares simulated and observed patterns
- treats validation as part of the model

Toy model or empirical model?

Toy ABM:

- asks whether a mechanism **could** generate a pattern
- emphasizes conceptual clarity
- uses stylized assumptions

Empirical ABM:

- constrains rules from data
- compares simulated and observed patterns
- treats validation as part of the model

Here we will build a toy model explicitly.

1. target phenomenon

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism
4. update schedule

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism
4. update schedule
5. outcome measure

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism
4. update schedule
5. outcome measure
6. sensitivity analysis

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism
4. update schedule
5. outcome measure
6. sensitivity analysis
7. comparison model

Design checklist for ABM studies

1. target phenomenon
2. agent-level mechanism
3. interaction mechanism
4. update schedule
5. outcome measure
6. sensitivity analysis
7. comparison model
8. validation target

Research question:

Research question:

Can population-level polarization emerge from simple local social influence?

Research question:

Can population-level polarization emerge from simple local social influence?

Belief state: b_i in $[-1, 1]$

Research question:

Can population-level polarization emerge from simple local social influence?

Belief state: b_i in $[-1, 1]$

- -1 : strong belief in position A

Research question:

Can population-level polarization emerge from simple local social influence?

Belief state: b_i in $[-1, 1]$

- -1 : strong belief in position A
- 0 : neutral

Research question:

Can population-level polarization emerge from simple local social influence?

Belief state: b_i in $[-1, 1]$

- -1 : strong belief in position A
- 0 : neutral
- $+1$: strong belief in position B

We use a ring network.

We use a ring network.

- each agent has two neighbors

We use a ring network.

- each agent has two neighbors
- interaction is local

We use a ring network.

- each agent has two neighbors
- interaction is local
- no global observer

We use a ring network.

- each agent has two neighbors
- interaction is local
- no global observer

We use a ring network.

- each agent has two neighbors
- interaction is local
- no global observer

The ring is deliberately unrealistic; it makes locality transparent.

0 -- 1 -- 2 -- 3 -- ... -- N-1
| |
+-----+



Local edges constrain who can influence whom.

Synchronous: all agents update simultaneously each step.

Synchronous: all agents update simultaneously each step.

Asynchronous: one agent is chosen per step (random or sequential).

Synchronous: all agents update simultaneously each step.

Asynchronous: one agent is chosen per step (random or sequential).

The choice changes dynamics and is often under-reported – it is part of the model specification.

At each step:

At each step:

1. choose one focal agent

At each step:

1. choose one focal agent
2. choose one neighbor

At each step:

1. choose one focal agent
2. choose one neighbor
3. compare beliefs

At each step:

1. choose one focal agent
2. choose one neighbor
3. compare beliefs
4. assimilate if similar

At each step:

1. choose one focal agent
2. choose one neighbor
3. compare beliefs
4. assimilate if similar
5. contrast if dissimilar

```
if  $|b_i - b_j| < \tau$ :  
     $b_i \leftarrow b_i + \alpha * (b_j - b_i)$   
else:  
     $b_i \leftarrow b_i - \beta * (b_j - b_i)$   
clip  $b_i$  to  $[-1, 1]$ 
```

- α : assimilation rate

- **alpha**: assimilation rate
- **beta**: contrast rate

- **alpha**: assimilation rate
- **beta**: contrast rate
- **tau**: tolerance threshold

- **alpha**: assimilation rate
- **beta**: contrast rate
- **tau**: tolerance threshold

- **alpha**: assimilation rate
- **beta**: contrast rate
- **tau**: tolerance threshold

These parameters define the cognitive-social mechanism.

Before coding, write the scientific contract

For this model, we need to specify:

Before coding, write the scientific contract

For this model, we need to specify:

- research question

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule
- outcome metrics

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule
- outcome metrics
- expected qualitative behavior

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule
- outcome metrics
- expected qualitative behavior
- verification checks

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule
- outcome metrics
- expected qualitative behavior
- verification checks

Before coding, write the scientific contract

For this model, we need to specify:

- research question
- agent state variables
- interaction topology
- update rule and schedule
- outcome metrics
- expected qualitative behavior
- verification checks

This is the same planning discipline we use with Cursor.

- initialize agents and beliefs

- initialize agents and beliefs
- define ring neighbors

- initialize agents and beliefs
- define ring neighbors
- implement `step()` update

- initialize agents and beliefs
- define ring neighbors
- implement `step()` update
- run simulation and record history

- initialize agents and beliefs
- define ring neighbors
- implement `step()` update
- run simulation and record history
- compute outcome measures

```
class BeliefPolarizationABM:
    def __init__(...): ...
    def neighbors(self, i): ...
    def step(self): ...
    def run(self, ...): ...
```

```
difference = self.beliefs[j] - self.beliefs[i]
if abs(difference) < self.tolerance:
    self.beliefs[i] += self.alpha * difference
else:
    self.beliefs[i] -= self.beta * difference
self.beliefs[i] = np.clip(self.beliefs[i], -1, 1)
```

- NetLogo: visual and educational; widely used in ABM

- NetLogo: visual and educational; widely used in ABM
- Mesa: Python ABM framework; integrates with NumPy, pandas, and matplotlib

- NetLogo: visual and educational; widely used in ABM
- Mesa: Python ABM framework; integrates with NumPy, pandas, and matplotlib
- AgentPy: lightweight Python option; useful for parameter sweeps

- NetLogo: visual and educational; widely used in ABM
- Mesa: Python ABM framework; integrates with NumPy, pandas, and matplotlib
- AgentPy: lightweight Python option; useful for parameter sweeps
- Plain Python: best here because every line of the mechanism stays visible

- mean belief

- mean belief
- variance

- mean belief
- variance
- polarization: mean $|b_i|$

- mean belief
- variance
- polarization: mean $|b_i|$
- extreme share: proportion near -1 or $+1$

- mean belief
- variance
- polarization: mean $|b_i|$
- extreme share: proportion near -1 or $+1$
- optional cluster count

Verification checks for this model

Useful checks include:

Useful checks include:

- beliefs always remain in $[-1, 1]$

Useful checks include:

- beliefs always remain in $[-1, 1]$
- the same seed reproduces the same trajectory

Useful checks include:

- beliefs always remain in $[-1, 1]$
- the same seed reproduces the same trajectory
- no-interaction or no-contrast variants behave as expected

Useful checks include:


- beliefs always remain in $[-1, 1]$
- the same seed reproduces the same trajectory
- no-interaction or no-contrast variants behave as expected
- output files are created in the expected folder

Verification checks for this model

Useful checks include:

- beliefs always remain in $[-1, 1]$
- the same seed reproduces the same trajectory
- no-interaction or no-contrast variants behave as expected
- output files are created in the expected folder
- parameter sweeps summarize multiple seeds, not one run

Individual trajectories



- agents start near the center


How to read trajectories

- agents start near the center
- local updates amplify differences


How to read trajectories

- agents start near the center
- local updates amplify differences
- repeated interaction can separate groups


Final distribution




Final distribution



Polarization over time



Polarization over time



What emerged?

- No agent tries to polarize the population.

What emerged?

- No agent tries to polarize the population.
- Each agent applies a local update rule.

What emerged?

- No agent tries to polarize the population.
- Each agent applies a local update rule.
- The population distribution changes shape because local updates accumulate over time.

ABMs often include stochastic elements:

ABMs often include stochastic elements:

- initial beliefs

ABMs often include stochastic elements:

- initial beliefs
- selected agents

ABMs often include stochastic elements:

- initial beliefs
- selected agents
- selected neighbors

ABMs often include stochastic elements:

- initial beliefs
- selected agents
- selected neighbors
- update order

ABMs often include stochastic elements:

- initial beliefs
- selected agents
- selected neighbors
- update order
- random seeds

ABMs often include stochastic elements:

- initial beliefs
- selected agents
- selected neighbors
- update order
- random seeds

ABMs often include stochastic elements:

- initial beliefs
- selected agents
- selected neighbors
- update order
- random seeds

Analyze the distribution over runs, not one trajectory.

Sensitivity analysis is required

One run is a demonstration, not evidence of a robust mechanism.

One run is a demonstration, not evidence of a robust mechanism.

- vary tolerance

Sensitivity analysis is required

One run is a demonstration, not evidence of a robust mechanism.

- vary tolerance
- vary contrast strength

Sensitivity analysis is required

One run is a demonstration, not evidence of a robust mechanism.


- vary tolerance
- vary contrast strength
- vary network structure

Sensitivity analysis is required

One run is a demonstration, not evidence of a robust mechanism.

- vary tolerance
- vary contrast strength
- vary network structure
- vary random seed

Parameter sweep over tolerance



- assimilation only (no contrast)

Comparison models to run

- assimilation only (no contrast)
- contrast only (no assimilation)

- assimilation only (no contrast)
- contrast only (no assimilation)
- random mixing (no network structure)

- assimilation only (no contrast)
- contrast only (no assimilation)
- random mixing (no network structure)
- ring vs small-world network

- assimilation only (no contrast)
- contrast only (no assimilation)
- random mixing (no network structure)
- ring vs small-world network
- heterogeneous tolerance

- assimilation only (no contrast)
- contrast only (no assimilation)
- random mixing (no network structure)
- ring vs small-world network
- heterogeneous tolerance
- external evidence source

- arbitrary rule choices

- arbitrary rule choices
- under-reported update schedule

- arbitrary rule choices
- under-reported update schedule
- sensitivity to implementation details

- arbitrary rule choices
- under-reported update schedule
- sensitivity to implementation details
- too many free parameters

- equifinality (different mechanisms produce the same macro-pattern)

- equifinality (different mechanisms produce the same macro-pattern)
- weak empirical grounding

- equifinality (different mechanisms produce the same macro-pattern)
- weak empirical grounding
- no serious baseline comparisons

- equifinality (different mechanisms produce the same macro-pattern)
- weak empirical grounding
- no serious baseline comparisons
- overclaiming generality from a toy model

What data constrain the model?

Empirical data should constrain:

What data constrain the model?

Empirical data should constrain:

- agent-level update rule

What data constrain the model?

Empirical data should constrain:

- agent-level update rule
- interaction network structure

What data constrain the model?

Empirical data should constrain:

- agent-level update rule
- interaction network structure
- time scale of updating

What data constrain the model?

Empirical data should constrain:

- agent-level update rule
- interaction network structure
- time scale of updating
- macro-pattern target

What would count as failure?

The model fails if:

What would count as failure?

The model fails if:

- polarization appears only for one random seed

What would count as failure?

The model fails if:

- polarization appears only for one random seed
- results vanish under plausible parameter ranges

What would count as failure?

The model fails if:

- polarization appears only for one random seed
- results vanish under plausible parameter ranges
- a simpler model produces the same pattern

What would count as failure?

The model fails if:

- polarization appears only for one random seed
- results vanish under plausible parameter ranges
- a simpler model produces the same pattern
- macro-patterns match but micro-dynamics are wrong

What would count as failure?

The model fails if:

- polarization appears only for one random seed
- results vanish under plausible parameter ranges
- a simpler model produces the same pattern
- macro-patterns match but micro-dynamics are wrong
- the update rule contradicts empirical behavioral data

ABM does not replace cognitive modeling.

ABM does not replace cognitive modeling.

ABM helps cognitive scientists ask whether a proposed cognitive mechanism, when embedded in an interaction structure, is **sufficient** to generate a population-level phenomenon.

You will build a small ABM that includes:

You will build a small ABM that includes:

- one research question with a testable hypothesis

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification
- explicit agent rules

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification
- explicit agent rules

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification
- explicit agent rules

Next week you will give a 5-10 minute presentation of your proposed project.

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification
- explicit agent rules

Next week you will give a 5-10 minute presentation of your proposed project.

We can all chime in on your project ideas and help you refine a research question.

You will build a small ABM that includes:

- one research question with a testable hypothesis
- a reproducible run script
- simulation-specific verification
- explicit agent rules

Next week you will give a 5-10 minute presentation of your proposed project.

We can all chime in on your project ideas and help you refine a research question.

ABM tests whether a proposed cognitive mechanism, embedded in an interaction structure, is sufficient to generate a target population-level phenomenon.

Axelrod, R. (1984). *The evolution of cooperation*. Basic Books.

Axelrod, R. (1984). *The evolution of cooperation*. Basic Books.

Epstein, J. M. (1999). Agent-based computational models and generative social science. *Complexity*, 4(5), 41-60. [https://doi.org/10.1002/\(SICI\)1099-0526\(199905/06\)4:5<41::AID-CPLX9>3.0.CO;2-F](https://doi.org/10.1002/(SICI)1099-0526(199905/06)4:5<41::AID-CPLX9>3.0.CO;2-F)

Axelrod, R. (1984). *The evolution of cooperation*. Basic Books.

Epstein, J. M. (1999). Agent-based computational models and generative social science. *Complexity*, 4(5), 41-60. [https://doi.org/10.1002/\(SICI\)1099-0526\(199905/06\)4:5<41::AID-CPLX9>3.0.CO;2-F](https://doi.org/10.1002/(SICI)1099-0526(199905/06)4:5<41::AID-CPLX9>3.0.CO;2-F)

Epstein, J. M., & Axtell, R. L. (1996). *Growing artificial societies: Social science from the bottom up*. The MIT Press. <https://doi.org/10.7551/mitpress/3374.001.0001>

Axelrod, R. (1984). *The evolution of cooperation*. Basic Books.

Epstein, J. M. (1999). Agent-based computational models and generative social science. *Complexity*, 4(5), 41-60. [https://doi.org/10.1002/\(SICI\)1099-0526\(199905/06\)4:5<41::AID-CPLX9>3.0.CO;2-F](https://doi.org/10.1002/(SICI)1099-0526(199905/06)4:5<41::AID-CPLX9>3.0.CO;2-F)

Epstein, J. M., & Axtell, R. L. (1996). *Growing artificial societies: Social science from the bottom up*. The MIT Press. <https://doi.org/10.7551/mitpress/3374.001.0001>

Schelling, T. C. (1971). Dynamic models of segregation. *Journal of Mathematical Sociology*, 1(2), 143-186. <https://doi.org/10.1080/0022250x.1971.9989794>